

## Assignment 4

Name: Ranjan Baro

ID: 2023PCP5274

### Program 1: Hello GPU

```
0s  %%writefile hello.cu
#include <stdio.h>
__global__ void hello() {
    printf("Hello, I am GPU.\n");
}
int main() {
    printf("Running Kernel...\n");
    hello<<<1,1>>>>();
    cudaDeviceSynchronize();
    printf("Hello, I am CPU\n");
    return 0;
}
```

Overwriting hello.cu

```
1s  [8] !nvcc hello.cu -o hello
```

```
0s  [9] !./hello
```

```
Running Kernel...
Hello, I am GPU.
Hello, I am CPU
```

### Program 2: Vector Addition

```
 %%writefile vector_add.cu
#include<iostream>
#define N (4*4)
#define THREADS_PER_BLOCK 4
using namespace std;
__global__ void addition(int *a, int *b, int *c, int *n) {
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    printf("(%d, %d)", index, *n);
    if(index < *n) {
        c[index] = a[index] + b[index];
    }
}
int main() {
    int *a, *b, *c, *d_a, *d_b, *d_c, *d_n;
    int size_int = sizeof(int); int size = N * sizeof(int); int n = N;
    cudaMalloc((void **)&d_a, size); cudaMalloc((void **)&d_b, size); cudaMalloc((void **)&d_c, size); cudaMalloc((void **)&d_n, size_int);
    a = (int *)malloc(size); b = (int *)malloc(size); c = (int *)malloc(size);
    for(int i = 0; i < N; i++) {
        a[i] = rand() % N; b[i] = rand() % N;
    }
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice); cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_n, &n, size_int, cudaMemcpyHostToDevice);
    addition<<<N/THREADS_PER_BLOCK, THREADS_PER_BLOCK>>>(d_a, d_b, d_c, d_n);
    cudaDeviceSynchronize(); cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost); printf("Addition of first 10 elements : \n");
    for(int i = 0; i < 10; i++) {
        printf("(%d + %d = %d), ", a[i], b[i], (c[i]));
    }
    free(a); free(b); free(c); cudaFree(a); cudaFree(b); cudaFree(c);
    return 0;
}
```

## Output:

```
Overwriting vector_add.cu

3s [!] !nvcc vector_add.cu -o add
!nvprof ./add

==22556== NVPROF is profiling process 22556, command: ./add
(8, 16)(9, 16)(10, 16)(11, 16)(0, 16)(1, 16)(2, 16)(3, 16)(4, 16)(5, 16)(6, 16)(7, 16)(12, 16)(13, 16)(14, 16)(15, 16)Additon of first 10 elements :
==22556== Profiling application: ./add
==22556== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 91.29% 72.735us 1 72.735us 72.735us 72.735us addition(int*, int*, int*, int*)
5.94% 4.7360us 3 1.5780us 1.3760us 1.9840us [CUDA memcpy HtoD]
2.77% 2.2080us 1 2.2080us 2.2080us 2.2080us [CUDA memcpy DtoH]
API calls: 99.69% 132.97ms 4 33.243ms 5.0810us 132.95ms cudaMalloc
0.09% 114.23us 101 1.1310us 130ns 47.118us cuDeviceGetAttribute
0.00% 113.02us 1 113.02us 113.02us 113.02us cudaDeviceSynchronize
0.06% 75.311us 4 18.827us 9.1340us 29.084us cudaMemcpy
0.05% 63.629us 1 63.629us 63.629us 63.629us cudaLaunchKernel
0.02% 28.384us 1 28.384us 28.384us 28.384us cuDeviceGetName
0.00% 5.8050us 1 5.8050us 5.8050us 5.8050us cuDeviceGetPCIBusId
0.00% 4.6300us 3 1.5430us 880ns 2.7270us cudaFree
0.00% 1.6130us 3 537ns 221ns 1.0870us cuDeviceGetCount
0.00% 1.1000us 2 550ns 312ns 788ns cuDeviceGet
0.00% 437ns 1 437ns 437ns 437ns cuDeviceTotalMem
0.00% 432ns 1 432ns 432ns 432ns cuModuleGetLoadingMode
0.00% 238ns 1 238ns 238ns 238ns cuDeviceGetUuid
(7 + 6 = 13), (9 + 3 = 12), (1 + 15 = 16), (10 + 12 = 22), (9 + 13 = 22), (10 + 11 = 21), (2 + 11 = 13), (3 + 6 = 9), (12 + 2 = 14), (4 + 8 = 12),
```

## Program 3: Matrix Multiplication

```
%%writefile matrix_multiplication.cu
#include <stdio.h>

__global__ void matrixMultKernel(int *a, int *b, int *c, int w) {
    int row = threadIdx.y + blockIdx.y * blockDim.y;
    int col = threadIdx.x + blockIdx.x * blockDim.x;
    int sum = 0;
    for(int k = 0; k < w; k++) {
        sum += a[row * w + k] * b[k * w + col];
    }
    c[row * w + col] = sum;
}

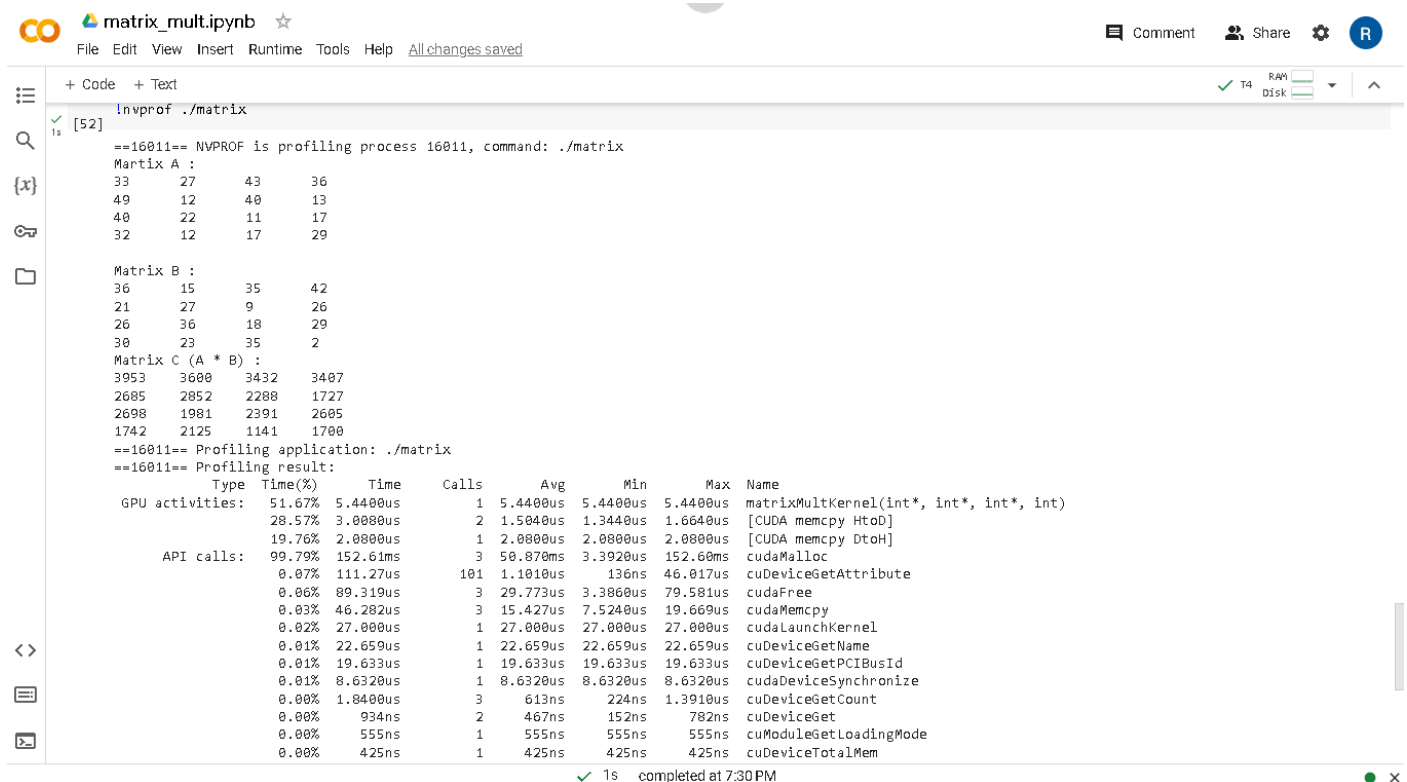
void matrixMult(int *a, int *b, int *c, int w) {
    int *d_a, *d_b, *d_c;
    int size = w * w * sizeof(int);
    cudaMalloc((void **) &d_a, size);
    cudaMalloc((void **) &d_b, size);
    cudaMalloc((void **) &d_c, size);
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
    dim3 blockDim(16, 16);
    dim3 gridDim((w + blockDim.x - 1) / blockDim.x, (w + blockDim.y - 1) / blockDim.y);
    matrixMultKernel<<<gridDim,blockDim>>>(d_a, d_b, d_c, w);
    cudaDeviceSynchronize();
    cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
}
```

```

int main() {
    const int w = 4;
    int a[w][w], b[w][w], c[w][w];
    for(int i = 0; i < w; i++) {
        for(int j = 0; j < w; j++) {
            a[i][j] = rand() % 50;
            b[i][j] = rand() % 50;
        }
    }
    matrixMult((int *)a, (int *)b, (int *)c, w);
    printf("Matrix A :\n");
    for(int i = 0; i < w; i++) {
        for(int j = 0; j < w; j++) {
            printf("%d\t", a[i][j]);
        }
        printf("\n");
    }
    printf("\nMatrix B :\n");
    for(int i = 0; i < w; i++) {
        for(int j = 0; j < w; j++) {
            printf("%d\t", b[i][j]);
        }
        printf("\n");
    }
    printf("Matrix C (A * B) :\n");
    for(int i = 0; i < w; i++) {
        for(int j = 0; j < w; j++) {
            printf("%d\t", c[i][j]);
        }
        printf("\n");
    }
    return 0;
}

```

## Output:



```
matrix_mult.ipynb
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
nvprof ./matrix
1s [52]
==16011== NVPROF is profiling process 16011, command: ./matrix
Matrix A :
33 27 43 36
49 12 40 13
40 22 11 17
32 12 17 29

Matrix B :
36 15 35 42
21 27 9 26
26 36 18 29
30 23 35 2

Matrix C (A * B) :
3953 3600 3432 3407
2685 2852 2288 1727
2698 1981 2391 2605
1742 2125 1141 1700
==16011== Profiling application: ./matrix
==16011== Profiling result:
Type Time(%) Time Calls Avg Min Max Name
GPU activities: 51.67% 5.4400us 1 5.4400us 5.4400us 5.4400us matrixMultKernel(int*, int*, int*, int)
28.57% 3.0000us 2 1.5000us 1.3440us 1.6640us [CUDA memcpy HtoD]
19.76% 2.0000us 1 2.0000us 2.0000us 2.0000us [CUDA memcpy DtoH]
API calls: 99.79% 152.61ms 3 50.870ms 3.3920us 152.60ms cudaMalloc
0.07% 111.27us 101 1.1010us 136ns 46.017us cuDeviceGetAttribute
0.06% 89.319us 3 29.773us 3.3860us 79.581us cudaFree
0.03% 46.282us 3 15.427us 7.5240us 19.660us cudaMemcpy
0.02% 27.000us 1 27.000us 27.000us 27.000us cudaLaunchKernel
0.01% 22.659us 1 22.659us 22.659us 22.659us cuDeviceGetName
0.01% 19.633us 1 19.633us 19.633us 19.633us cuDeviceGetPCIBusId
0.01% 8.6320us 1 8.6320us 8.6320us 8.6320us cuDeviceSynchronize
0.00% 1.0400us 3 613ns 224ns 1.3910us cuDeviceGetCount
0.00% 934ns 2 467ns 152ns 782ns cuDeviceGet
0.00% 555ns 1 555ns 555ns 555ns cuModuleGetLoadingMode
0.00% 425ns 1 425ns 425ns 425ns cuDeviceTotalMem
1s completed at 7:30 PM
```

## Summary Report On CUDA:

**CUDA:** CUDA, which stands for Compute Unified Device Architecture, is a parallel computing platform and programming model that makes using a GPU for general purpose computing simple.

### ➤ Keywords:

- **Host – CPU**
- **Device – GPU**
- **Host Memory – System memory (DRAM) associated with host**
- **Device Memory – GPU memory**
- **Kernel – Function executed on GPU by single thread**

### Key components of CUDA:

- GPU Utilization:** CUDA leverages the parallel processing capabilities of NVIDIA GPUs, allowing developers to perform general-purpose computations in parallel, leading to significant speedups.
- CUDA Toolkit:** NVIDIA provides a software development kit called the CUDA Toolkit, which includes the necessary tools, libraries, and programming environments for GPU development.
- Programming Languages:** CUDA supports programming in CUDA C and CUDA C++, extensions of the C and C++ languages, enabling developers to write parallel programs explicitly controlling GPU execution.
- Runtime API:** The CUDA Runtime API provides a set of functions that C/C++ programs can call to manage and control the execution of GPU kernels.
- Parallel Architecture:** NVIDIA GPUs are designed with multiple cores, allowing the simultaneous execution of numerous parallel threads. CUDA takes advantage of this architecture to accelerate a wide range of applications, including scientific simulations, deep learning, and image processing.

## **CUDA Memory**

### **a) Global Memory:**

- Global memory is the largest and slowest memory in the CUDA hierarchy.
- It is accessible by all threads on the GPU and is used to store global variables.

### **b) Shared Memory:**

- Shared by all threads in a thread block
- Faster and small
- used for communication and data sharing between threads within the same block.

### **c) Local Memory:**

- slower than registers and is typically used when a thread's register usage is high.
- Per thread private memory

### **d) Texture Memory:**

- a read-only cache designed for texture fetching.
- optimized for 2D spatial locality.
- can be beneficial for certain types of memory access patterns.